

---

# DS203 : EXERCISE 3

---

**Nirav Bhattad**

Last updated September 1, 2024

## Step 1

### Step 1

Review the Jupyter Notebook `E3.ipynb` and:

- (a) Create a summary of the code therein.
- (b) Are there any learnings from this code that you wish to highlight?

This Python Notebook `E3.ipynb` presents an analysis of various regression models applied to a dataset using polynomial features of different degrees. The models analyzed include Linear Regression, Support Vector Machine (SVM) Regression, Random Forest, Gradient Boosting, K-Nearest Neighbors (KNN), and Neural Networks. The performance of each model is evaluated based on metrics such as R-squared, Mean Squared Error (MSE), Durbin-Watson, and Jarque-Bera statistics.

We use the dataset given in `E3-MLR3.xlsx`. The dataset contains 548 samples of  $(y, x_1)$  to train the model and 152 samples of  $(y, x_1)$  to test the model.

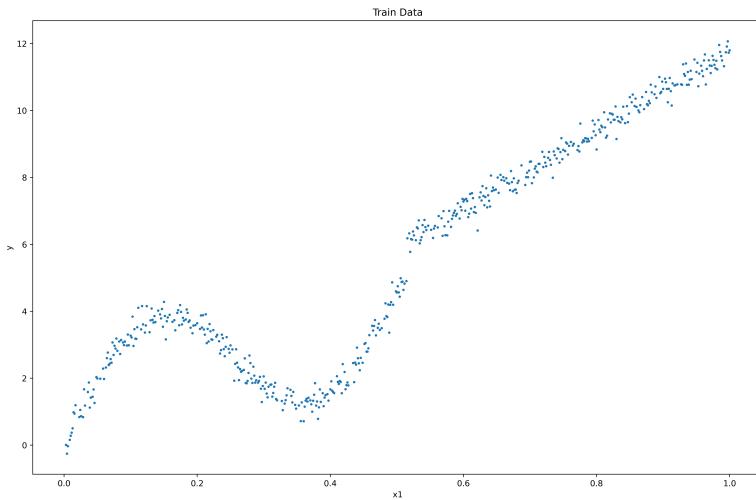


Figure 1: Training Data used in the analysis

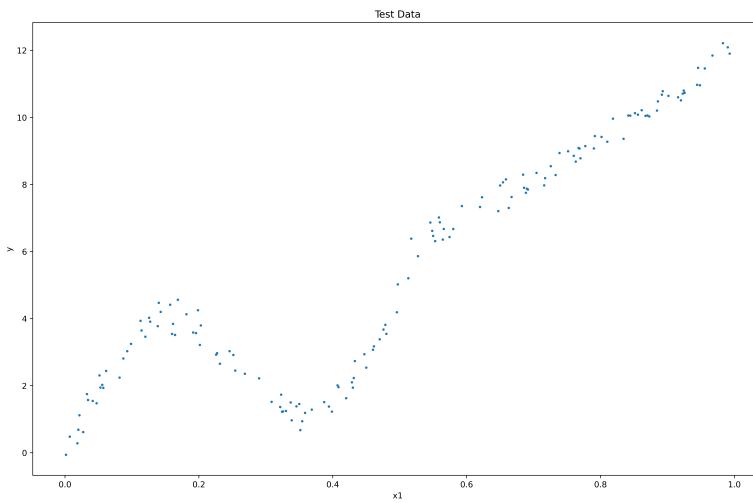


Figure 2: Testing Data used in the analysis

The following Python script in `E3.ipynb` applies multiple regression algorithms to a dataset, evaluates their performance, and visualizes the results.

## 1. Importing Necessary Libraries

```

1 import pandas as pd
2 import numpy as np
3 from sklearn.preprocessing import PolynomialFeatures
4 from sklearn.linear_model import LinearRegression
5 from sklearn.svm import SVR
6 from sklearn.ensemble import RandomForestRegressor
7 from sklearn.ensemble import GradientBoostingRegressor
8 from sklearn.neural_network import MLPRegressor
9 from sklearn.neighbors import KNeighborsRegressor
10 from sklearn.metrics import mean_squared_error, r2_score
11 import statsmodels.api as sm
12 import matplotlib.pyplot as plt

```

Listing 1: Importing Libraries

**Description:** The code begins by importing essential libraries for data manipulation (`pandas`, `numpy`), machine learning models (`scikit-learn`), statistical analysis (`statsmodels`), and plotting (`matplotlib`).

## 2. Data Loading and Preprocessing

```

1 file_path = 'E3-MLR3.xlsx'
2 train_data = pd.read_excel(file_path, sheet_name='train')
3 test_data = pd.read_excel(file_path, sheet_name='test')
4
5 X_train = train_data.drop(columns=['y'])
6 y_train = train_data['y']
7
8 X_test = test_data.drop(columns=['y'])
9 y_test = test_data['y']
10
11 poly = PolynomialFeatures(degree=6)
12 X_train_poly = poly.fit_transform(X_train)

```

```
13 X_test_poly = poly.transform(X_test)
```

Listing 2: Loading and Preprocessing Data

**Description:** The script reads training and testing data from an Excel file and separates features (`X_train`, `X_test`) and the target variable (`y_train`, `y_test`). It then uses `PolynomialFeatures` to augment the features with polynomial terms of degree 6.

### 3. Saving the Augmented Dataset

```
1 feature_names = poly.get_feature_names_out(X_train.columns)
2 augmented_data = pd.DataFrame(X_test_poly, columns=feature_names)
3 augmented_data['y'] = test_data['y']
4 augmented_data.to_csv('augmented_data.csv', index=False)
```

Listing 3: Saving Augmented Data

**Description:** The code generates a CSV file containing the test data with the newly created polynomial features for review.

### 4. Defining and Running Regression Models

```
1 algorithms = {
2     'Linear Regression': LinearRegression(),
3     'SVM Regression': SVR(kernel='poly'),
4     'RandomForest': RandomForestRegressor(),
5     'XGBoost': GradientBoostingRegressor(),
6     'knn': KNeighborsRegressor(),
7     'Neural Network': MLPRegressor(hidden_layer_sizes=[10,10,10], max_iter=20000)
8 }
9
10 metric_table_train = pd.DataFrame()
11 metric_table_test = pd.DataFrame()
12
13 fig, axs = plt.subplots(len(algorithms), 4, figsize=(20, 20))
14 fig_row = -1
15
16 for algorithm_name, algorithm in algorithms.items():
17     algorithm.fit(X_train_poly, y_train)
18     y_train_pred = algorithm.predict(X_train_poly)
19     y_test_pred = algorithm.predict(X_test_poly)
20
21     r2_train = algorithm.score(X_train_poly, y_train)
22     mse_train = mean_squared_error(y_train, y_train_pred)
23     r2_test = algorithm.score(X_test_poly, y_test)
24     mse_test = mean_squared_error(y_test, y_test_pred)
25
26     residuals_train = y_train - y_train_pred
27     residuals_test = y_test - y_test_pred
28
29     durbin_watson_stat_train = sm.stats.durbin_watson(residuals_train)
30     jb_stat_train, jb_p_value_train, _, _ = sm.stats.jarque_bera(residuals_train)
31
32     durbin_watson_stat_test = sm.stats.durbin_watson(residuals_test)
33     jb_stat_test, jb_p_value_test, _, _ = sm.stats.jarque_bera(residuals_test)
34
35     metric_table_train.at[algorithm_name, 'R-squared'] = r2_train
36     metric_table_train.at[algorithm_name, 'MSE'] = mse_train
37     metric_table_train.at[algorithm_name, 'Durbin-Watson'] = durbin_watson_stat_train
38     metric_table_train.at[algorithm_name, 'Jarque-Bera'] = jb_stat_train
39     metric_table_train.at[algorithm_name, 'JB P-value'] = jb_p_value_train
40
41     metric_table_test.at[algorithm_name, 'R-squared'] = r2_test
42     metric_table_test.at[algorithm_name, 'MSE'] = mse_test
43     metric_table_test.at[algorithm_name, 'Durbin-Watson'] = durbin_watson_stat_test
```

```
44 metric_table_test.at[algorithm_name, 'Jarque-Bera'] = jb_stat_test
45 metric_table_test.at[algorithm_name, 'JB P-value'] = jb_p_value_test
46
47 fig_row = fig_row+1
48
49 axs[fig_row, 0].scatter(train_data['x1'], y_train)
50 axs[fig_row, 0].scatter(train_data['x1'], y_train_pred)
51 axs[fig_row, 0].set_title(algorithm_name + " - Train")
52
53 axs[fig_row, 1].scatter(train_data['x1'], residuals_train)
54 axs[fig_row, 1].set_title(algorithm_name + " Residuals - Train")
55
56 axs[fig_row, 2].scatter(test_data['x1'], y_test)
57 axs[fig_row, 2].scatter(test_data['x1'], y_test_pred)
58 axs[fig_row, 2].set_title(algorithm_name + " - Test")
59
60 axs[fig_row, 3].scatter(test_data['x1'], residuals_test)
61 axs[fig_row, 3].set_title(algorithm_name + " Residuals - Test")
```

Listing 4: Defining and Running Regression Models

**Description:** A dictionary of regression algorithms is created. Each algorithm is trained on the polynomial features, and its performance is evaluated using metrics such as R-squared, Mean Squared Error (MSE), Durbin-Watson, and Jarque-Bera test. These metrics are stored in tables. The residuals and predicted values are also plotted.

## 5. Visualizing the Results

```
1 plt.tight_layout()
2 plt.show()
```

Listing 5: Visualizing the Results

**Description:** The code concludes with displaying the generated plots for each regression model.

### Learnings from this Code

- (a) The code demonstrates the application of various regression algorithms to a dataset with polynomial features.
- (b) It highlights the importance of evaluating regression models using multiple metrics to understand their performance.
- (c) The code showcases the use of residual analysis to assess the model's assumptions and identify potential issues.
- (d) It provides a structured approach to comparing different regression models and their outcomes.
- (e) The code emphasizes the significance of visualizing the model's predictions and residuals for better interpretation.
- (f) The code also shows how to save augmented datasets for further analysis or review.

## Step 2

### Step 2

Review the **Sklearn** documentation for each Sklearn function used in the Notebook (eg. **PolynomialFeatures**, **LinearRegression**, **mean\_squared\_error**, etc.) and create a description of each to explain, to yourself, the functionality, the input parameters, and the outputs generated. Present this in the form of a two-column - Table (Function name | Description).

<b>PolynomialFeatures</b>	<p>This function is used to generate polynomial and interaction features. It generates a new feature matrix consisting of all polynomial combinations of the features with a degree less than or equal to the specified degree. The input to this function is the degree of the polynomial features to be generated. The output generated is a new feature matrix consisting of all polynomial combinations of the features with a degree less than or equal to the specified degree.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>degree</code>: int or tuple (<code>min_degree</code>, <code>max_degree</code>), default=2</li> <li>• <code>interaction_only</code>: bool, default=False</li> <li>• <code>include_bias</code>: bool, default=True</li> <li>• <code>order</code>: str in {'C', 'F'}, default='C'</li> </ul> <p><b>Attributes:</b></p> <ul style="list-style-type: none"> <li>• <code>powers_</code>: ndarray of shape (<code>n_output_features_</code>, <code>n_input_features_</code>)</li> <li>• <code>n_output_features_</code>: int</li> <li>• <code>n_features_in_</code>: int</li> <li>• <code>feature_names_in_</code>: ndarray of shape (<code>n_input_features_</code>,)</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• ndarray of shape (<code>n_samples</code>, <code>n_output_features_</code>)</li> </ul>
<b>LinearRegression</b>	<p>This function is used to fit a linear model. It fits a linear model with coefficients <math>w = (w_1, \dots, w_p)</math> to minimize the residual sum of squares between the observed targets in the dataset and the targets predicted by the linear approximation.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>fit_intercept</code>: bool, default=True</li> <li>• <code>copy_X</code>: bool, default=True</li> <li>• <code>n_jobs</code>: int, default=None</li> <li>• <code>positive</code>: bool, default=False</li> </ul> <p><b>Attributes:</b></p> <ul style="list-style-type: none"> <li>• <code>coef_</code>: ndarray of shape (<code>n_targets</code>, <code>n_features</code>)</li> <li>• <code>intercept_</code>: ndarray of shape (<code>n_targets</code>,)</li> <li>• <code>rank_</code>: int</li> <li>• <code>singular_</code>: ndarray of shape (<code>min(X, y)</code>,)</li> <li>• <code>n_features_in_</code>: ndarray of shape (<code>n_targets</code>,)</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• <code>self</code>: returns an instance of self</li> </ul>

SVR	<p>This function is used to fit the Support Vector Regression model.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>kernel</code>: str, default='rbf'</li> <li>• <code>degree</code>: int, default=3</li> <li>• <code>gamma</code>: float, default='scale'</li> <li>• <code>coef0</code>: float, default=0.0</li> <li>• <code>tol</code>: float, default=1e-3</li> <li>• <code>C</code>: float, default=1.0</li> <li>• <code>epsilon</code>: float, default=0.1</li> <li>• <code>shrinking</code>: bool, default=True</li> <li>• <code>cache_size</code>: float, default=200</li> <li>• <code>verbose</code>: bool, default=False</li> <li>• <code>max_iter</code>: int, default=-1</li> </ul> <p><b>Attributes:</b></p> <ul style="list-style-type: none"> <li>• <code>coef_</code>: ndarray of shape (1, n_features)</li> <li>• <code>dual_coef_</code>: ndarray of shape (1, n_SV)</li> <li>• <code>fit_status_</code>: int</li> <li>• <code>intercept_</code>: ndarray of shape (1,)</li> <li>• <code>n_features_in_</code>: int</li> <li>• <code>feature_names_in_</code>: ndarray of shape (n_features,)</li> <li>• <code>n_iter_</code>: int</li> <li>• <code>n_support_</code>: ndarray of shape (1,)</li> <li>• <code>shape_fit_</code>: tuple of int of shape (n_dimensions_of_X,)</li> <li>• <code>support_</code>: ndarray of shape (n_SV,)</li> <li>• <code>support_vectors_</code>: ndarray of shape (n_SV, n_features)</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• ndarray of shape (n_samples,)</li> </ul>
KNeighborsRegressor	<p>This function is used to fit the K-Nearest Neighbors regressor model, which is a non-parametric method used for regression. It predicts the target variable by learning from the <math>k</math> nearest data points in the feature space.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>n_neighbors</code>: int, default=5</li> <li>• <code>weights</code>: str, default='uniform'</li> <li>• <code>algorithm</code>: str, default='auto'</li> <li>• <code>leaf_size</code>: int, default=30</li> <li>• <code>p</code>: int, default=2</li> <li>• <code>metric</code>: str, default='minkowski'</li> <li>• <code>metric_params</code>: dict, default=None</li> <li>• <code>n_jobs</code>: int, default=None</li> </ul> <p><b>Attributes:</b></p> <ul style="list-style-type: none"> <li>• <code>effective_metric_</code>: str</li> <li>• <code>effective_metric_params_</code>: dict</li> <li>• <code>n_samples_fit_</code>: int</li> <li>• <code>n_features_in_</code>: int</li> <li>• <code>feature_names_in_</code>: ndarray of shape (n_features,)</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• ndarray of shape (n_samples,)</li> </ul>

<b>RandomForestRegressor</b>	<p>This function is used to fit a random forest model, which is an ensemble learning method for regression. It fits a number of classifying decision trees on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>n_estimators</code>: int, default=100</li> <li>• <code>criterion</code>: str, default='squared_error'</li> <li>• <code>max_depth</code>: int, default=None</li> <li>• <code>min_samples_split</code>: int, default=2</li> <li>• <code>min_samples_leaf</code>: int, default=1</li> <li>• <code>min_weight_fraction_leaf</code>: float, default=0.0</li> <li>• <code>max_features</code>: int, default='auto'</li> <li>• <code>max_leaf_nodes</code>: int, default=None</li> <li>• <code>min_impurity_decrease</code>: float, default=0.0</li> <li>• <code>bootstrap</code>: bool, default=True</li> <li>• <code>oob_score</code>: bool, default=False</li> <li>• <code>n_jobs</code>: int, default=None</li> <li>• <code>random_state</code>: int, default=None</li> <li>• <code>verbose</code>: int, default=0</li> <li>• <code>warm_start</code>: bool, default=False</li> <li>• <code>ccp_alpha</code>: float, default=0.0</li> <li>• <code>max_samples</code>: int, default=None</li> <li>• <code>monotonic_cst</code>: array-like of shape (n_features), default=None</li> </ul> <p><b>Attributes:</b></p> <ul style="list-style-type: none"> <li>• <code>estimator_</code>: DecisionTreeRegressor</li> <li>• <code>estimators_</code>: list of DecisionTreeRegressor</li> <li>• <code>n_features_in_</code>: int</li> <li>• <code>feature_names_in_</code>: ndarray of shape (n_features,)</li> <li>• <code>n_outputs_</code>: int</li> <li>• <code>oob_score_</code>: ndarray of shape (n_estimators,)</li> <li>• <code>oob_prediction_</code>: ndarray of shape (n_samples, n_outputs)</li> <li>• <code>estimators_samples_</code>: list of ndarray</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• ndarray of shape (n_samples,)</li> </ul>
<b>mean_squared_error</b>	<p>This function is used to compute the mean squared error regression loss. It computes the mean squared error between the true and predicted values.</p> <p><b>Input Parameters:</b></p> <ul style="list-style-type: none"> <li>• <code>y_true</code>: ndarray of shape (n_samples,)</li> <li>• <code>y_pred</code>: ndarray of shape (n_samples,)</li> <li>• <code>squared</code>: bool, default=True</li> <li>• <code>multioutput</code>: str, default='uniform_average'</li> <li>• <code>sample_weight</code>: ndarray of shape (n_samples,), default=None</li> </ul> <p><b>Output:</b></p> <ul style="list-style-type: none"> <li>• float</li> </ul>

<b>GradientBoostingRegressor</b>	<p>This function is used to fit a gradient boosting model, which is an ensemble learning method for regression. It fits a number of regression trees on various sub-samples of the dataset and uses boosting to improve the predictive accuracy and control over-fitting.</p> <p>Input Parameters:</p> <ul style="list-style-type: none"><li>• <code>loss</code>: str, default='ls'</li><li>• <code>learning_rate</code>: float, default=0.1</li><li>• <code>n_estimators</code>: int, default=100</li><li>• <code>subsample</code>: float, default=1.0</li><li>• <code>criterion</code>: str, default='friedman_mse'</li><li>• <code>min_samples_split</code>: int, default=2</li><li>• <code>min_samples_leaf</code>: int, default=1</li><li>• <code>min_weight_fraction_leaf</code>: float, default=0.0</li><li>• <code>max_depth</code>: int, default=3</li><li>• <code>min_impurity_decrease</code>: float, default=0.0</li><li>• <code>init</code>: estimator, default=None</li><li>• <code>random_state</code>: int, default=None</li><li>• <code>max_features</code>: str, default=None</li><li>• <code>alpha</code>: float, default=0.9</li><li>• <code>verbose</code>: int, default=0</li><li>• <code>max_leaf_nodes</code>: int, default=None</li><li>• <code>warm_start</code>: bool, default=False</li><li>• <code>validation_fraction</code>: float, default=0.1</li><li>• <code>n_iter_no_change</code>: int, default=None</li><li>• <code>tol</code>: float, default=1e-4</li><li>• <code>ccp_alpha</code>: float, default=0.0</li></ul> <p>Attributes:</p> <ul style="list-style-type: none"><li>• <code>n_estimators_</code>: int</li><li>• <code>n_trees_per_iteration_</code>: int</li><li>• <code>feature_importances_</code>: ndarray of shape (n_features,)</li><li>• <code>oob_improvement_</code>: ndarray of shape (n_estimators,)</li><li>• <code>oob_score_</code>: ndarray of shape (n_estimators,)</li><li>• <code>oob_scores_</code>: ndarray of shape (n_estimators,)</li><li>• <code>train_score_</code>: ndarray of shape (n_estimators,)</li><li>• <code>init_</code>: estimator</li><li>• <code>estimators_</code>: ndarray of DecisionTreeRegressor</li><li>• <code>n_features_in_</code>: int</li><li>• <code>feature_names_in_</code>: ndarray of shape (n_features,)</li><li>• <code>max_features_</code>: int</li></ul> <p>Output:</p> <ul style="list-style-type: none"><li>• ndarray of shape (n_samples,)</li></ul>
----------------------------------	---

<b>r2_score</b>	This function is used to compute the R-squared regression score function. It computes the R-squared score, which is the coefficient of determination. This function has not been used in the jupyter notebook. <b>Input Parameters:</b> <ul style="list-style-type: none"><li>• <b>y_true</b>: ndarray of shape (n_samples,)</li><li>• <b>y_pred</b>: ndarray of shape (n_samples,)</li><li>• <b>sample_weight</b>: ndarray of shape (n_samples,), default=None</li><li>• <b>multioutput</b>: str, default='uniform_average'</li><li>• <b>force_finite</b>: bool, default=True</li></ul> <b>Output:</b> <ul style="list-style-type: none"><li>• <b>float</b></li></ul>
-----------------	---

<b>MLPRegressor</b>	<p>This function is used to fit a Multi-layer Perceptron regressor. It trains a neural network model with a single hidden layer, which is a fully-connected feed-forward artificial neural network.</p> <p>Input Parameters:</p> <ul style="list-style-type: none"><li>• <code>hidden_layer_sizes</code>: tuple, default=(100,)</li><li>• <code>activation</code>: str, default='relu'</li><li>• <code>solver</code>: str, default='adam'</li><li>• <code>alpha</code>: float, default=0.0001</li><li>• <code>batch_size</code>: int, default='auto'</li><li>• <code>learning_rate</code>: str, default='constant'</li><li>• <code>learning_rate_init</code>: float, default=0.001</li><li>• <code>power_t</code>: float, default=0.5</li><li>• <code>max_iter</code>: int, default=200</li><li>• <code>shuffle</code>: bool, default=True</li><li>• <code>random_state</code>: int, default=None</li><li>• <code>tol</code>: float, default=1e-4</li><li>• <code>verbose</code>: bool, default=False</li><li>• <code>warm_start</code>: bool, default=False</li><li>• <code>momentum</code>: float, default=0.9</li><li>• <code>nesterovs_momentum</code>: bool, default=True</li><li>• <code>early_stopping</code>: bool, default=False</li><li>• <code>validation_fraction</code>: float, default=0.1</li><li>• <code>beta_1</code>: float, default=0.9</li><li>• <code>beta_2</code>: float, default=0.999</li><li>• <code>epsilon</code>: float, default=1e-8</li><li>• <code>n_iter_no_change</code>: int, default=10</li><li>• <code>max_fun</code>: int, default=15000</li></ul> <p>Attributes:</p> <ul style="list-style-type: none"><li>• <code>loss_</code>: float</li><li>• <code>best_loss_</code>: float</li><li>• <code>loss_curve_</code>: list of shape (n_iter_,)</li><li>• <code>t_</code>: int</li><li>• <code>validation_scores_</code>: list of shape (n_iter_,)</li><li>• <code>best_validation_score_</code>: float</li><li>• <code>coefs_</code>: list of shape (n_layers_ - 1,)</li><li>• <code>intercepts_</code>: list of shape (n_layers_ - 1,)</li><li>• <code>n_features_in_</code>: int</li><li>• <code>feature_names_in_</code>: ndarray of shape (n_features,)</li><li>• <code>n_iter_</code>: int</li><li>• <code>n_layers_</code>: int</li><li>• <code>n_outputs_</code>: int</li><li>• <code>out_activation_</code>: str</li></ul> <p>Output:</p> <ul style="list-style-type: none"><li>• ndarray of shape (n_samples,)</li></ul>
---------------------	--

## Step 3

### Step 3

Generate outputs by setting **degree = 1, degree = 3, degree = 6, degree = 10**, in the `PolynomialFeatures` function used in `E3.ipynb` and analyze them as follows:

- (a) Review the `augmented_data.csv` file generated in each case and document your observations.
- (b) Create an overall qualitative summary based on a review and analysis of the Figures generated.
- (c) Summarize and explain the variations in the metrics **across regression methods for a given degree** (ie. a given set of polynomial features). Cover both, train and test, metrics, and compare them.
- (d) Summarize and explain the variations in the metrics **across degrees for a given regression method**. Cover both, train, and test metrics, and compare them.
- (e) When **degree = 1** which method(s) result in acceptable regression models? Why?
- (f) When **degree = 6** which method(s) result in acceptable regression models? Why?
- (g) As the value of degree is increased to 10 which regression methods show the most impact? Why?
- (h) Why do non-parametric methods like KNN and `Decision Tree` based methods generate good results even without feature engineering?
- (i) What are the limitations of the non-parametric methods?
- (j) Given the results, should `LinearRegression` be used at all? Why, when? Justify your answer.

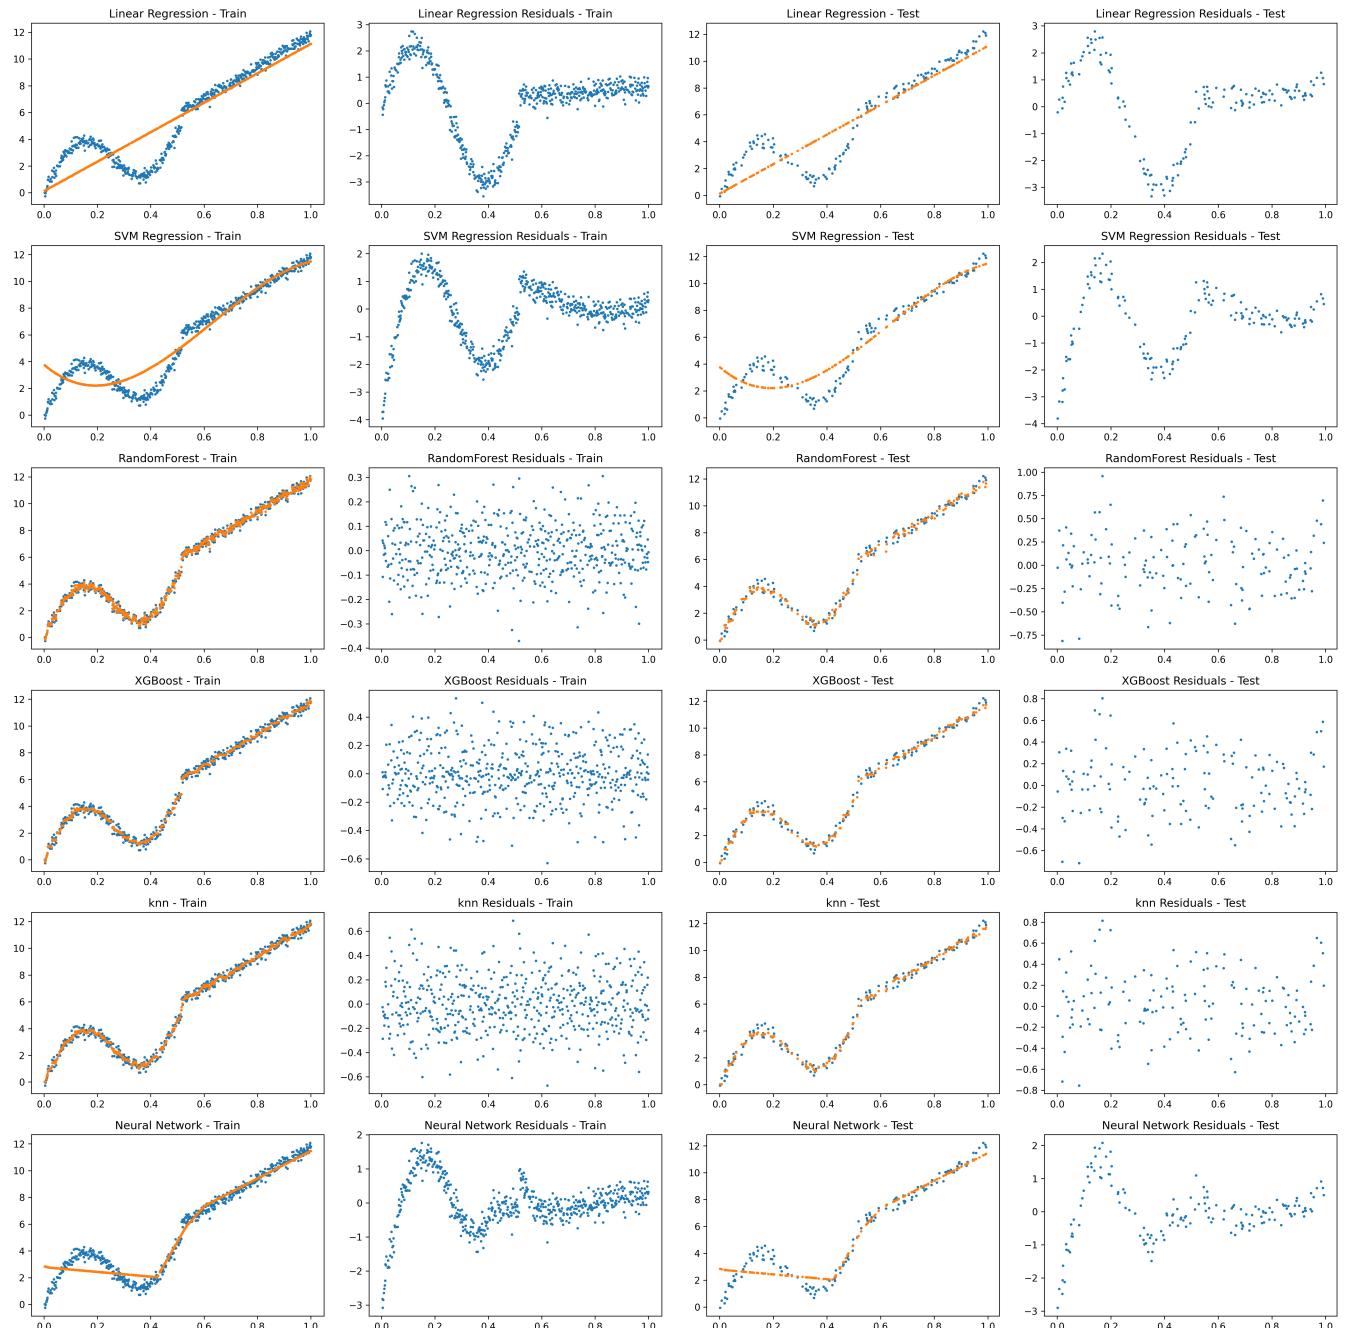


Figure 3: Various Regression Models with degree = 1

When degree = 1, the following observations can be made:

-

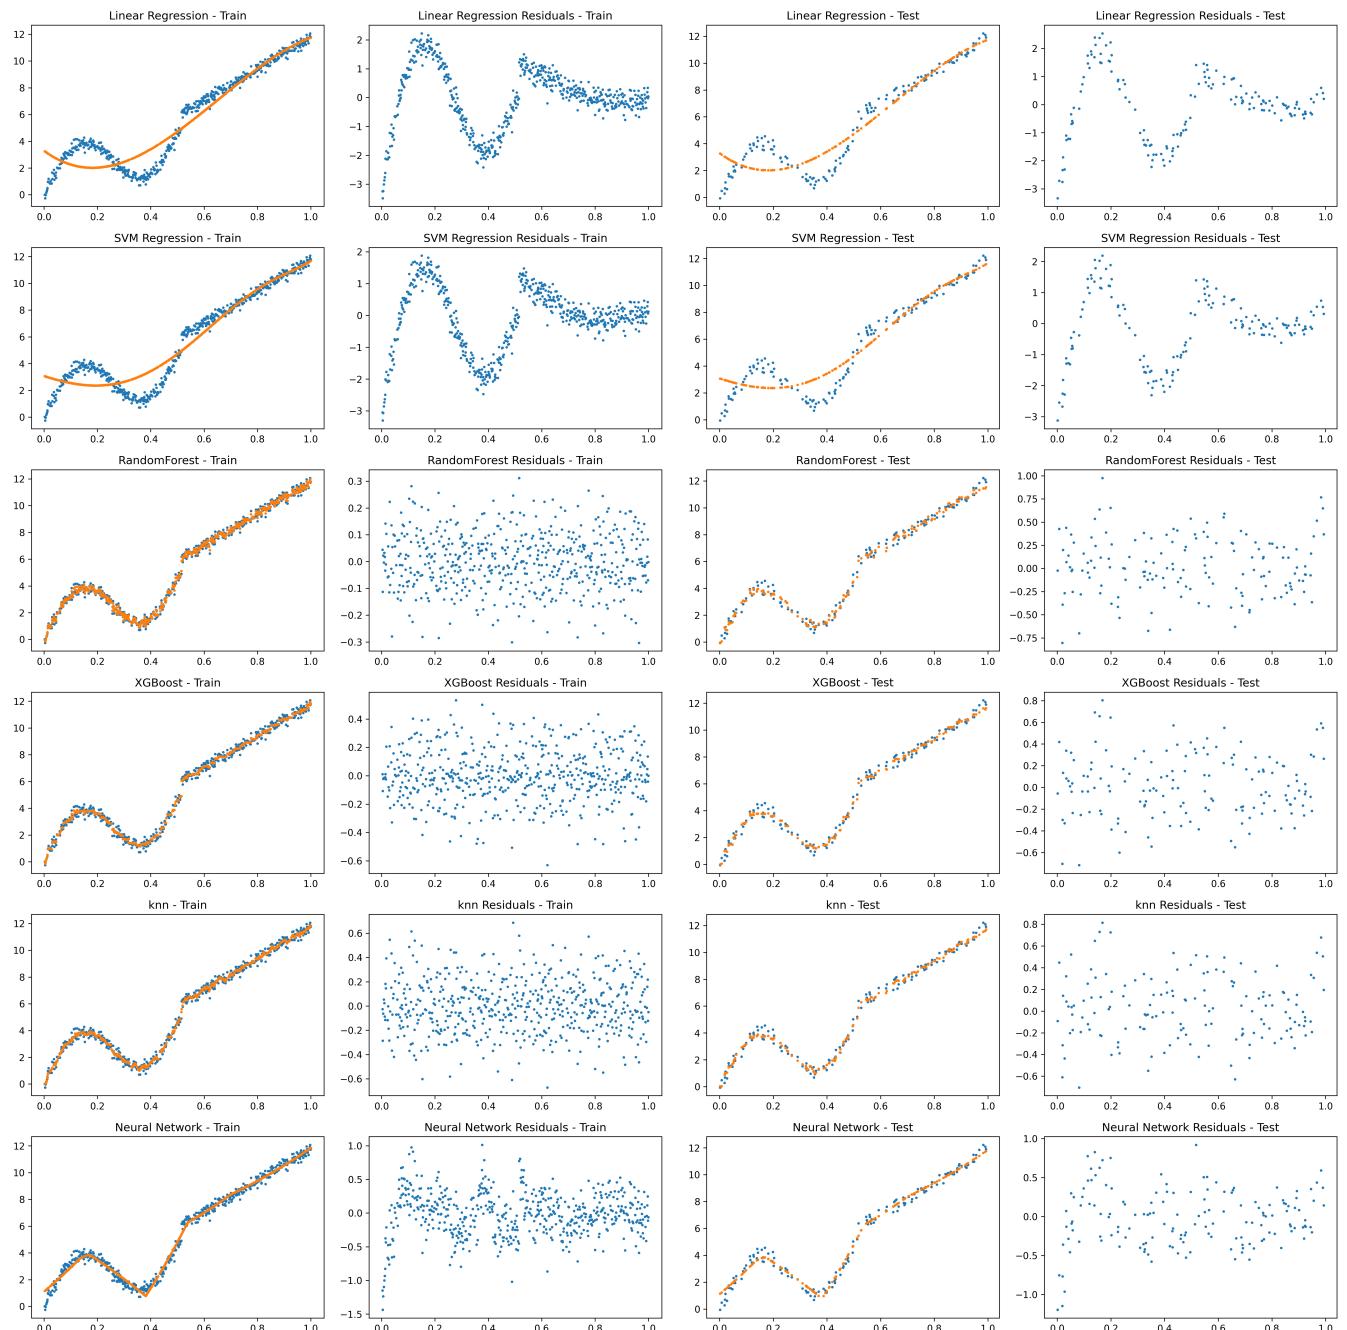


Figure 4: Various Regression Models with degree = 3

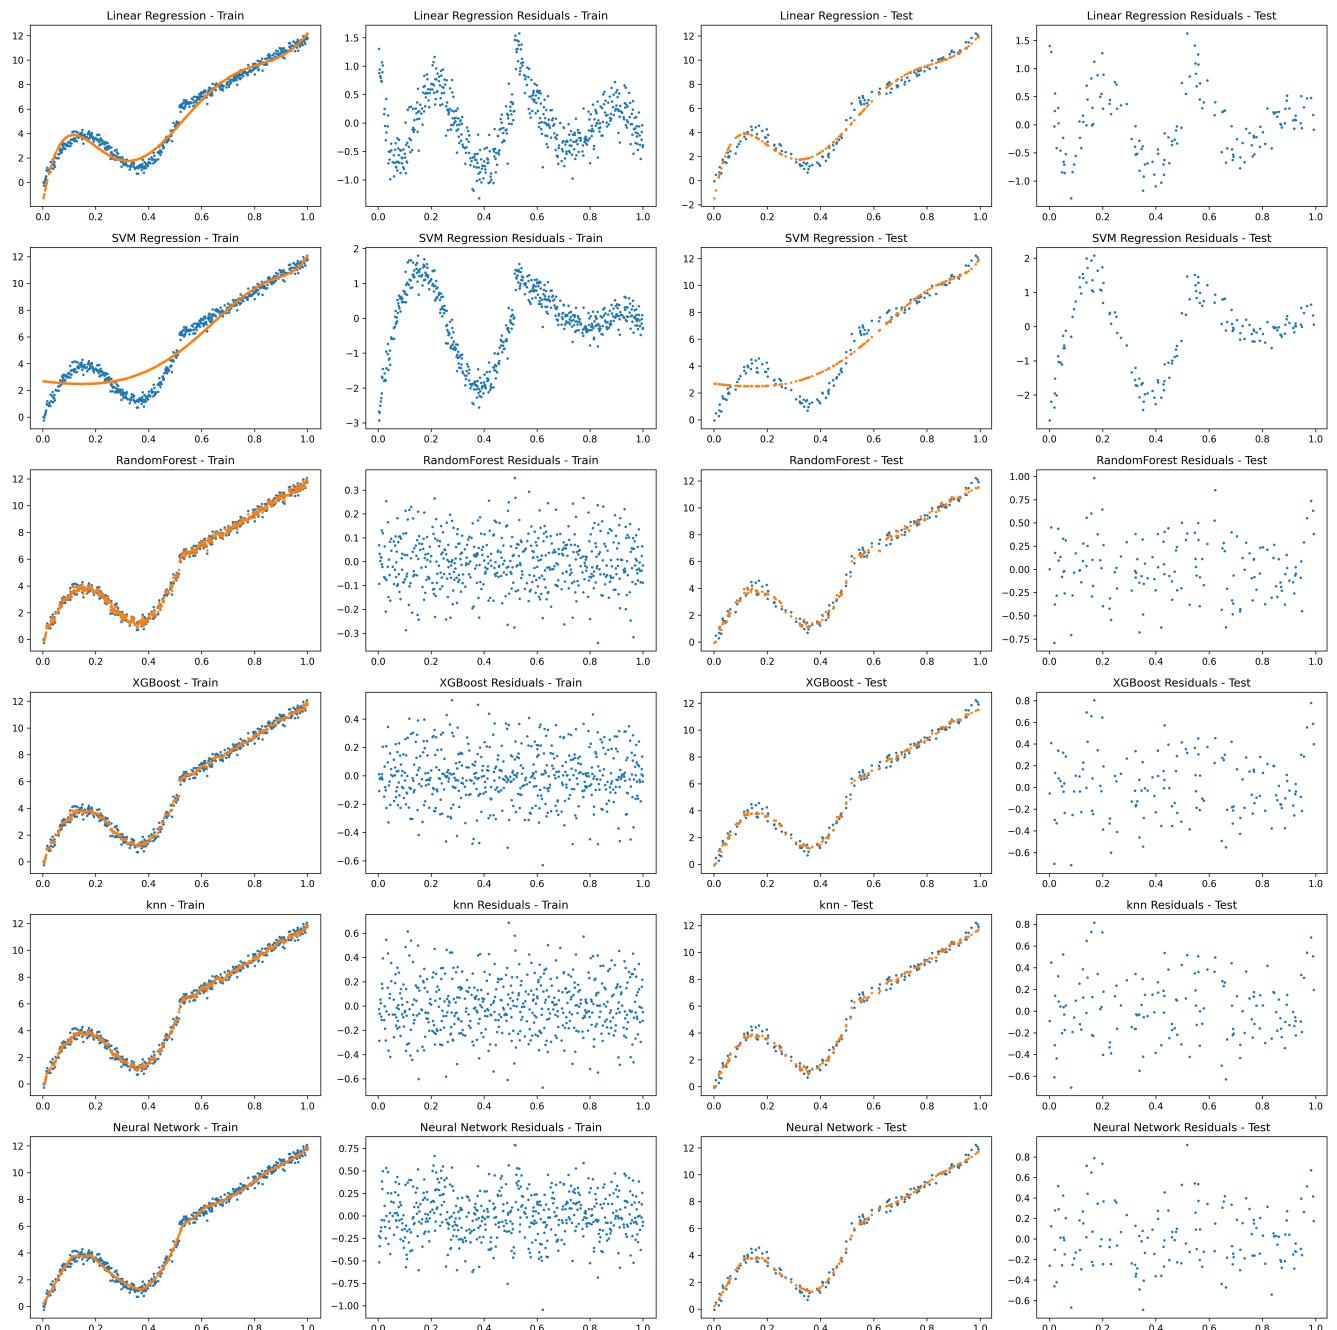


Figure 5: Various Regression Models with degree = 6

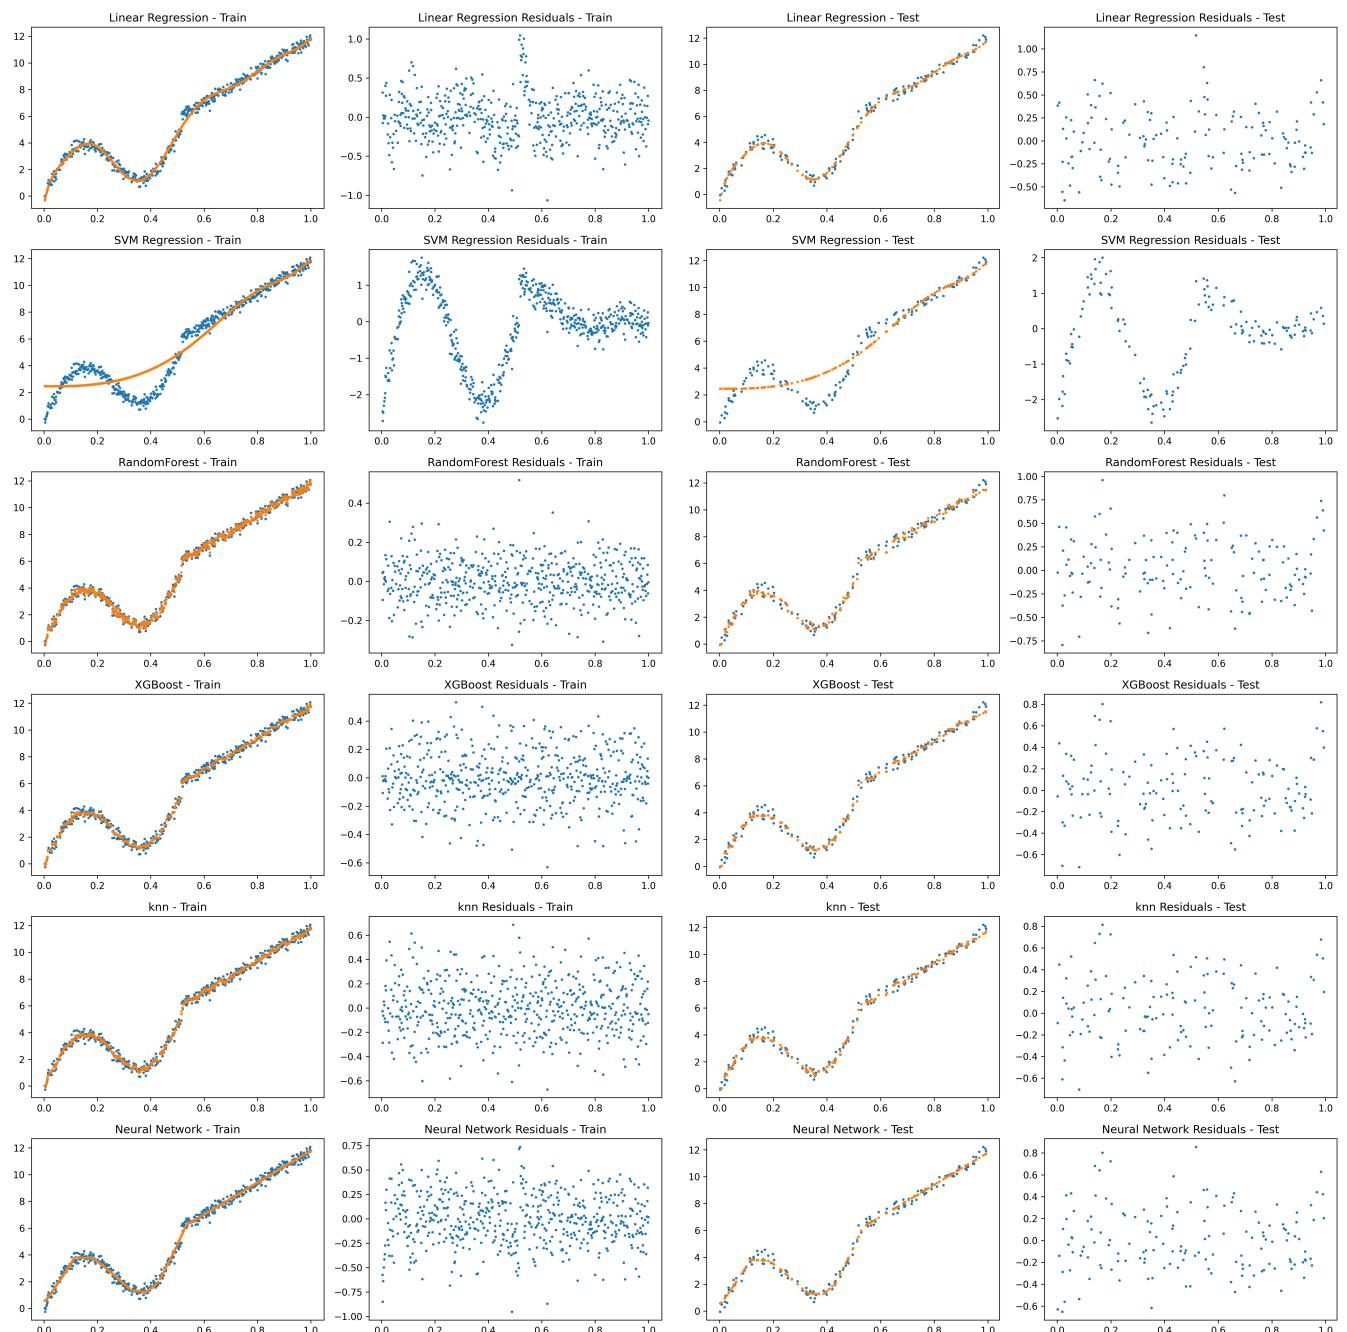


Figure 6: Various Regression Models with degree = 10

**Question (a)**

Review the `augmented_data.csv` file generated in each case and document your observations.

**Question (b)**

Create an overall qualitative summary based on a review and analysis of the Figures generated.

**Question (c)**

Summarize and explain the variations in the metrics **across regression methods for a given degree** (ie. a given set of polynomial features). Cover both, train and test, metrics, and compare them.

**Question (d)**

Summarize and explain the variations in the metrics **across degrees for a given regression method**. Cover both, train, and test metrics, and compare them.

**Question (e)**

When **degree = 1** which method(s) result in acceptable regression models? Why?

**Question (f)**

When **degree = 6** which method(s) result in acceptable regression models? Why?

**Question (g)**

As the value of degree is increased to 10 which regression methods show the most impact? Why?

**Question (h)**

Why do non-parametric methods like KNN and **Decision Tree** based methods generate good results even without feature engineering?

**Question (i)**

What are the limitations of the non-parametric methods?

**Question (j)**

Given the results, should LinearRegression be used at all? Why, when? Justify your answer.

## Step 4

### Step 4

In step 2 you have already reviewed the important parameters and outputs related to the regression methods. Select 2-3 methods, vary the important parameters, and observe how the outputs change (eg. see the function calls for `SVR` and `MLPRegressor`). Document the outcomes of your experiments.

## Step 5

### Step 5

Review **Sklearn** documentation to understand and experiment with a few more (2-3) regression methods, in addition to the ones listed above, and document the outcomes of your experiments.

## Main Learnings